

```

import numpy as np
import pandas as pd
#Visualization
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score

```

1. Problem Statement
#To predict price of cars using various features

2. Data Gathering
df = pd.read_csv("/content/autos_dataset.csv")
df

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	?	alfa-romero	gas	std	two	convertible	rwd	fro
1	3	?	alfa-romero	gas	std	two	convertible	rwd	fro
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	fro
3	2	164	audi	gas	std	four	sedan	fwd	fro
4	2	164	audi	gas	std	four	sedan	4wd	fro
...
200	-1	95	volvo	gas	std	four	sedan	rwd	fro
201	-1	95	volvo	gas	turbo	four	sedan	rwd	fro
202	-1	95	volvo	gas	std	four	sedan	rwd	fro
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	fro
204	-1	95	volvo	gas	turbo	four	sedan	rwd	fro

205 rows × 26 columns

df.head() # head() function is generally use for to print first 5 rows by default

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136

5 rows × 26 columns

df.head().T # after use of T after head() function, Columns goes to rows place and Rows goes to Columns place

	0	1	2	3	4	
symboling	3	3	1	2	2	
normalized-losses	?	?	?	164	164	
make	alfa-romero	alfa-romero	alfa-romero	audi	audi	
fuel-type	gas	gas	gas	gas	gas	
aspiration	std	std	std	std	std	
num-of-doors	two	two	two	four	four	
body-style	convertible	convertible	hatchback	sedan	sedan	
drive-wheels	rwd	rwd	rwd	fwd	4wd	
engine-location	front	front	front	front	front	
wheel-base	88.6	88.6	94.5	99.8	99.4	
length	168.8	168.8	171.2	176.6	176.6	
width	64.1	64.1	65.5	66.2	66.4	
height	48.8	48.8	52.4	54.3	54.3	
curb-weight	2548	2548	2823	2337	2824	
engine-type	dohc	dohc	ohcv	ohc	ohc	
num-of-cylinders	four	four	six	four	five	
engine-size	130	130	152	109	136	
fuel-system	mpfi	mpfi	mpfi	mpfi	mpfi	
bore	3.47	3.47	2.68	3.19	3.19	
stroke	2.68	2.68	3.47	3.4	3.4	
compression-ratio	9.0	9.0	9.0	10.0	8.0	
horsepower	111	111	154	102	115	

```
# 3. Exploratory Data Analysis (EDA) :
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   symboling        205 non-null    int64  
 1   normalized-losses 205 non-null    object  
 2   make              205 non-null    object  
 3   fuel-type         205 non-null    object  
 4   aspiration        205 non-null    object  
 5   num-of-doors      205 non-null    object  
 6   body-style        205 non-null    object  
 7   drive-wheels      205 non-null    object  
 8   engine-location    205 non-null    object  
 9   wheel-base         205 non-null    float64
 10  length             205 non-null    float64
 11  width              205 non-null    float64
 12  height             205 non-null    float64
 13  curb-weight        205 non-null    int64  
 14  engine-type        205 non-null    object  
 15  num-of-cylinders   205 non-null    object  
 16  engine-size         205 non-null    int64  
 17  fuel-system         205 non-null    object  
 18  bore                205 non-null    object  
 19  stroke              205 non-null    object  
 20  compression-ratio   205 non-null    float64
 21  horsepower          205 non-null    object  
 22  peak-rpm            205 non-null    object  
 23  city-mpg            205 non-null    int64  
 24  highway-mpg          205 non-null    int64  
 25  price               205 non-null    object  
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB
```

```
df.isna().sum() # Checking there is avl. of any NaN values
```

```
symboling      0
normalized-losses 0
make          0
fuel-type     0
aspiration    0
num-of-doors  0
body-style    0
drive-wheels  0
engine-location 0
wheel-base    0
length        0
width         0
height         0
curb-weight   0
engine-type   0
num-of-cylinders 0
engine-size   0
fuel-system   0
bore          0
stroke         0
compression-ratio 0
horsepower    0
peak-rpm       0
city-mpg       0
highway-mpg    0
price          0
dtype: int64
```

"""Check all the columns one by one, fill the NaN values , Convert "object" Data Types into "int" or "float" bcz during Model Training and Model valuation Calculation , it is must requirement of integers"""

'Check all the columns one by one, fill the NaN values , Convert "object" Data Types in to "int" or "float" bcz during\nModel Training and Model valuation Calculation , it is must requirement of integers'

```
# 3.1 Symboling
```

```
df[\"symboling\"]
```

0	3
1	3
2	1
3	2
4	2
..	
200	-1
201	-1
202	-1
203	-1
204	-1

```
Name: symboling, Length: 205, dtype: int64
```

```
# 3.2 normalized-losses
```

```
df[\"normalized-losses\"]
```

0	?
1	?
2	?
3	164
4	164
..	
200	95
201	95
202	95
203	95
204	95

```
Name: normalized-losses, Length: 205, dtype: object
```

""In the return output we see question marks and its data type is "object" , so this ? question mark indicate it is NaN Value.
It is necessary to convert this ? question marks to NaN standard format.
Because it is generate some confusion because its data type is "object" and at isna().sum() function there is also shows
there is not available of NaN values.
So we replace ? Question mark into NaN Value."""

'In the return output we see question marks and its data type is "object" , so this ? question mark indicate it is NaN Value.\nIt is necessary to convert this ? question marks to NaN standard format.\nBecause it is generate some confusion because its data type is "object" and at isna().sum() function there is also shows\nthere is not available of Na

```
df.replace({"?":np.nan},inplace=True) # By using replace function we replace all question mark with NaN Values
```

```
df.isna().sum() # After replace of ? question mark we find summary of availabilty of NaN Values.
```

symboling	0
normalized-losses	41
make	0
fuel-type	0
aspiration	0
num-of-doors	2
body-style	0
drive-wheels	0
engine-location	0
wheel-base	0
length	0
width	0
height	0
curb-weight	0
engine-type	0
num-of-cylinders	0
engine-size	0
fuel-system	0
bore	4
stroke	4
compression-ratio	0
horsepower	2
peak-rpm	2
city-mpg	0
highway-mpg	0
price	4
	dtype: int64

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>			
RangeIndex: 205 entries, 0 to 204			
Data columns (total 26 columns):			
#	Column	Non-Null Count	Dtype
0	symboling	205 non-null	int64
1	normalized-losses	164 non-null	object
2	make	205 non-null	object
3	fuel-type	205 non-null	object
4	aspiration	205 non-null	object
5	num-of-doors	203 non-null	object
6	body-style	205 non-null	object
7	drive-wheels	205 non-null	object
8	engine-location	205 non-null	object
9	wheel-base	205 non-null	float64
10	length	205 non-null	float64
11	width	205 non-null	float64
12	height	205 non-null	float64
13	curb-weight	205 non-null	int64
14	engine-type	205 non-null	object
15	num-of-cylinders	205 non-null	object
16	engine-size	205 non-null	int64
17	fuel-system	205 non-null	object
18	bore	201 non-null	object
19	stroke	201 non-null	object
20	compression-ratio	205 non-null	float64
21	horsepower	203 non-null	object
22	peak-rpm	203 non-null	object
23	city-mpg	205 non-null	int64
24	highway-mpg	205 non-null	int64
25	price	201 non-null	object

dtypes: float64(5), int64(5), object(16)

memory usage: 41.8+ KB

```
df["normalized-losses"] = df["normalized-losses"].astype(float) # "normalized-losses" this column in "object" data type
# So we converted into "float" data type
```

```
df["normalized-losses"].mean()
```

```
122.0
```

```
df["normalized-losses"].median()
```

```
115.0
```

**"""We check both Mean and Median for "normalized-losses" column but Mean value is slighter greater than Median.
This is because of availability of some Outliers.
So we use Median to fill the NaN Values."""**

**'We check both Mean and Median for "normalized-losses" column but Mean value is slighte
r greater than Median.\nThis is because of availability of some Outliers.\nSo we use Med
ian to fill the NaN Values.'**

```
df["normalized-losses"] = df["normalized-losses"].fillna(df["normalized-losses"].median()).astype(int)
# we fill the NaN Values by Median of "normalized-losses" column and converted into "int" data type
```

```
df["normalized-losses"]
```

```
0      115
1      115
2      115
3      164
4      164
...
200     95
201     95
202     95
203     95
204     95
Name: normalized-losses, Length: 205, dtype: int64
```

```
df.info() # By doing Type Casting we convert "object" into "int" data type and
# We fill the NaN Values by Median of "normalized-losses" column
# Like wise that we process all Columns( Fill NaN Values and Converted into "int" or "float")
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   symboling        205 non-null    int64  
 1   normalized-losses 205 non-null    int64  
 2   make             205 non-null    object  
 3   fuel-type        205 non-null    object  
 4   aspiration       205 non-null    object  
 5   num-of-doors     203 non-null    object  
 6   body-style       205 non-null    object  
 7   drive-wheels     205 non-null    object  
 8   engine-location   205 non-null    object  
 9   wheel-base       205 non-null    float64 
 10  length           205 non-null    float64 
 11  width            205 non-null    float64 
 12  height           205 non-null    float64 
 13  curb-weight      205 non-null    int64  
 14  engine-type      205 non-null    object  
 15  num-of-cylinders 205 non-null    object  
 16  engine-size      205 non-null    int64  
 17  fuel-system      205 non-null    object  
 18  bore              201 non-null    object  
 19  stroke            201 non-null    object  
 20  compression-ratio 205 non-null    float64 
 21  horsepower        203 non-null    object  
 22  peak-rpm          203 non-null    object  
 23  city-mpg          205 non-null    int64  
 24  highway-mpg        205 non-null    int64  
 25  price             201 non-null    object  
dtypes: float64(5), int64(6), object(15)
memory usage: 41.8+ KB
```

```
# 3.3 make
df["make"]
```

```

0      alfa-romero
1      alfa-romero
2      alfa-romero
3      audi
4      audi
...
200     volvo
201     volvo
202     volvo
203     volvo
204     volvo
Name: make, Length: 205, dtype: object

```

```
df["make"].nunique() # We find out how many type of unique values availale in " make " column
```

```
22
```

```
df["make"].value_counts() # We find out how many type of unique values availale and its Count
```

make	Count
toyota	32
nissan	18
mazda	17
mitsubishi	13
honda	13
volkswagen	12
subaru	12
peugot	11
volvo	11
dodge	9
mercedes-benz	8
bmw	8
audi	7
plymouth	7
saab	6
porsche	5
isuzu	4
jaguar	3
chevrolet	3
alfa-romero	3
renault	2
mercury	1

```
Name: make, dtype: int64
```

```
#3.4 fuel-type
```

```
df["fuel-type"].value_counts()
```

fuel-type	Count
gas	185
diesel	20

```
Name: fuel-type, dtype: int64
```

```
df["fuel-type"].replace({'gas' : 0 , 'diesel' : 1}, inplace = True)
```

```
df["fuel-type"].value_counts()
```

fuel-type	Count
0	185
1	20

```
Name: fuel-type, dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   symboling        205 non-null    int64  
 1   normalized-losses 205 non-null    int64  
 2   make             205 non-null    object  
 3   fuel-type        205 non-null    int64  
 4   aspiration       205 non-null    object  
 5   num-of-doors     203 non-null    object  
 6   body-style       205 non-null    object  
 7   drive-wheels     205 non-null    object  

```

```

8   engine-location    205 non-null    object
9   wheel-base        205 non-null    float64
10  length            205 non-null    float64
11  width             205 non-null    float64
12  height            205 non-null    float64
13  curb-weight       205 non-null    int64
14  engine-type       205 non-null    object
15  num-of-cylinders  205 non-null    object
16  engine-size       205 non-null    int64
17  fuel-system       205 non-null    object
18  bore              201 non-null    object
19  stroke             201 non-null    object
20  compression-ratio 205 non-null    float64
21  horsepower         203 non-null    object
22  peak-rpm           203 non-null    object
23  city-mpg           205 non-null    int64
24  highway-mpg        205 non-null    int64
25  price              201 non-null    object
dtypes: float64(5), int64(7), object(14)
memory usage: 41.8+ KB

```

3.5 aspiration

```
df['aspiration'].value_counts()
```

```

std      168
turbo     37
Name: aspiration, dtype: int64

```

```
df['aspiration'].replace({'std' : 0 , 'turbo' : 1}, inplace = True)
```

```
df['aspiration'].value_counts()
```

```

0      168
1      37
Name: aspiration, dtype: int64

```

3.6 num-of-doors

```
df['num-of-doors'].value_counts()
```

```

four    114
two     89
Name: num-of-doors, dtype: int64

```

```
df['num-of-doors'].replace({'four' : 4 , 'two' : 2}, inplace = True)
df['num-of-doors'].value_counts()
```

```

4.0    114
2.0    89
Name: num-of-doors, dtype: int64

```

```
df['num-of-doors'].mean()
```

```
3.123152709359606
```

```
df['num-of-doors'].median()
```

```
4.0
```

```
df['num-of-doors'].mode()[0]
```

```
4.0
```

```
df['num-of-doors'].fillna(df['num-of-doors'].median() , inplace = True)
```

##3.7 body-style

```
df['body-style'].value_counts()
```

```

sedan      96
hatchback 70
wagon     25
hardtop     8
convertible 6
Name: body-style, dtype: int64

```

df.T

	0	1	2	3	4	5	6	7	8	9	...	195	196	197	198
symboling	3	3	1	2	2	2	1	1	1	0	...	-1	-2	-1	-2
normalized-losses	115	115	115	164	164	115	158	115	158	115	...	74	103	74	103
make	alfa-romero	alfa-romero	alfa-romero	audi	audi	audi	audi	audi	audi	audi	...	volvo	volvo	volvo	volvo
fuel-type	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0
aspiration	0	0	0	0	0	0	0	0	1	1	...	0	0	0	1
num-of-doors	2.0	2.0	2.0	4.0	4.0	2.0	4.0	4.0	4.0	2.0	...	4.0	4.0	4.0	4.0
body-style	convertible	convertible	hatchback	sedan	sedan	sedan	wagon	sedan	hatchback	...	wagon	sedan	wagon	sedan	
drive-wheels	rwd	rwd	rwd	fwd	4wd	fwd	fwd	fwd	fwd	4wd	...	rwd	rwd	rwd	rwd
engine-location	front	front	front	front	front	front	front	front	front	front	...	front	front	front	front
wheel-base	88.6	88.6	94.5	99.8	99.4	99.8	105.8	105.8	105.8	99.5	...	104.3	104.3	104.3	104.3
length	168.8	168.8	171.2	176.6	176.6	177.3	192.7	192.7	192.7	178.2	...	188.8	188.8	188.8	188.8
width	64.1	64.1	65.5	66.2	66.4	66.3	71.4	71.4	71.4	67.9	...	67.2	67.2	67.2	67.2
height	48.8	48.8	52.4	54.3	54.3	53.1	55.7	55.7	55.9	52.0	...	57.5	56.2	57.5	56.2
curb-weight	2548	2548	2823	2337	2824	2507	2844	2954	3086	3053	...	3034	2935	3042	3045
engine-type	dohc	dohc	ohcv	ohc	ohc	ohc	ohc	ohc	ohc	ohc	...	ohc	ohc	ohc	ohc
num-of-cylinders	four	four	six	four	five	five	five	five	five	five	...	four	four	four	four
engine-size	130	130	152	109	136	136	136	136	131	131	...	141	141	141	130
fuel-system	mpfi	mpfi	mpfi	mpfi	mpfi	mpfi	mpfi	mpfi	mpfi	mpfi	...	mpfi	mpfi	mpfi	mpfi
bore	3.47	3.47	2.68	3.19	3.19	3.19	3.19	3.19	3.13	3.13	...	3.78	3.78	3.78	3.62
stroke	2.68	2.68	3.47	3.4	3.4	3.4	3.4	3.4	3.4	3.4	...	3.15	3.15	3.15	3.15
compression-ratio	9.0	9.0	9.0	10.0	8.0	8.5	8.5	8.5	8.3	7.0	...	9.5	9.5	9.5	7.5
horsepower	111	111	154	102	115	110	110	110	140	160	...	114	114	114	162
peak-rpm	5000	5000	5000	5500	5500	5500	5500	5500	5500	5500	...	5400	5400	5400	5100
city-mpg	21	21	19	24	18	19	19	19	17	16	...	23	24	24	17
highway-mpg	27	27	26	30	22	25	25	25	20	22	...	28	28	28	22
price	13495	16500	16500	13950	17450	15250	17710	18920	23875	NaN	...	13415	15985	16515	18420

26 rows × 205 columns

```

df = pd.get_dummies(df,columns = ['body-style'])
df

```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	...	horsepower	peak-rpm	city-mpg	highway-mpg
0	3	115	alfa-romero	0	0	2.0	rwd	front	88.6	168.8	...	111	5000	21	27
1	3	115	alfa-romero	0	0	2.0	rwd	front	88.6	168.8	...	111	5000	21	27
2	1	115	alfa-romero	0	0	2.0	rwd	front	94.5	171.2	...	154	5000	19	26
3	2	164	audi	0	0	4.0	fwd	front	99.8	176.6	...	102	5500	24	30
4	2	164	audi	0	0	4.0	4wd	front	99.4	176.6	...	115	5500	18	22
...

#3.8 drive-wheels

df['drive-wheels'].value_counts()

```
fwd    120
rwd    76
4wd     9
Name: drive-wheels, dtype: int64
```

df['drive-wheels'].replace({"fwd" : 0 , "rwd" : 1 , "4wd" : 2}, inplace=True)
df['drive-wheels'].value_counts()

```
0    120
1    76
2     9
Name: drive-wheels, dtype: int64
```

3.9 engine-location#

df['engine-location'].value_counts()

```
front    202
rear      3
Name: engine-location, dtype: int64
```

df['engine-location'].replace({"front" : 0 , "rear" : 1 }, inplace=True)
df['engine-location'].value_counts()

```
0    202
1      3
Name: engine-location, dtype: int64
```

3.14 engine-type

df['engine-type'].value_counts()

```
ohc    148
ohcf    15
ohcv    13
dohc    12
l      12
rotor     4
dohcv     1
Name: engine-type, dtype: int64
```

df = pd.get_dummies(df, columns=['engine-type'])
df

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	...	style	body-style_hatchback	body-style_sedan	body-style_wagon
0	3	115	alfa-romero	0	0	2.0	1	0	88.6	168.8	0	0	0
1	3	115	alfa-romero	0	0	2.0	1	0	88.6	168.8	0	0	0
2	1	115	alfa-romero	0	0	2.0	1	0	94.5	171.2	1	0	0
3	2	164	audi	0	0	4.0	0	0	99.8	176.6	0	1	0
4	2	164	audi	0	0	4.0	2	0	99.4	176.6	0	1	0
...
200	-1	95	volvo	0	0	4.0	1	0	109.1	188.8	0	1	0

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 36 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   symboling        205 non-null    int64  
 1   normalized-losses 205 non-null    int64  
 2   make             205 non-null    object  
 3   fuel-type        205 non-null    int64  
 4   aspiration       205 non-null    int64  
 5   num-of-doors     205 non-null    float64 
 6   drive-wheels     205 non-null    int64  
 7   engine-location   205 non-null    int64  
 8   wheel-base       205 non-null    float64 
 9   length            205 non-null    float64 
 10  width             205 non-null    float64 
 11  height            205 non-null    float64 
 12  curb-weight      205 non-null    int64  
 13  num-of-cylinders 205 non-null    object  
 14  engine-size       205 non-null    int64  
 15  fuel-system       205 non-null    object  
 16  bore              201 non-null    object  
 17  stroke            201 non-null    object  
 18  compression-ratio 205 non-null    float64 
 19  horsepower         203 non-null    object  
 20  peak-rpm           203 non-null    object  
 21  city-mpg          205 non-null    int64  
 22  highway-mpg        205 non-null    int64  
 23  price              201 non-null    object  
 24  body-style_convertible 205 non-null    uint8  
 25  body-style_hardtop  205 non-null    uint8  
 26  body-style_hatchback 205 non-null    uint8  
 27  body-style_sedan   205 non-null    uint8  
 28  body-style_wagon   205 non-null    uint8  
 29  engine-type_dohc   205 non-null    uint8  
 30  engine-type_dohcv  205 non-null    uint8  
 31  engine-type_l       205 non-null    uint8  
 32  engine-type_ohc    205 non-null    uint8  
 33  engine-type_ohcf   205 non-null    uint8  
 34  engine-type_ohcv   205 non-null    uint8  
 35  engine-type_rotor  205 non-null    uint8  
dtypes: float64(6), int64(10), object(8), uint8(12)
memory usage: 41.0+ KB
```

3.15 num-of-cylinders

df['num-of-cylinders'].value_counts()

four	159
six	24
five	11
eight	5
two	4
three	1
twelve	1

Name: num-of-cylinders, dtype: int64

df['num-of-cylinders'].replace({'four' : 4 , 'six' : 6 , 'five' : 5 , 'eight' : 8 , 'two' : 2 , 'three' : 3,'twelve' : 12}, inplace = True)

```
df['num-of-cylinders'].value_counts()
```

4	159
6	24
5	11
8	5
2	4
3	1
12	1

Name: num-of-cylinders, dtype: int64

```
df
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	...	style_hatchback	body-style_sedan	body-style_suv
0	3	115	alfa-romero	0	0	2.0	1	0	88.6	168.8	...	0	0	0
1	3	115	alfa-romero	0	0	2.0	1	0	88.6	168.8	...	0	0	0
2	1	115	alfa-romero	0	0	2.0	1	0	94.5	171.2	...	1	0	0
3	2	164	audi	0	0	4.0	0	0	99.8	176.6	...	0	1	0
4	2	164	audi	0	0	4.0	2	0	99.4	176.6	...	0	1	0
...
200	-1	95	volvo	0	0	4.0	1	0	109.1	188.8	...	0	1	0
201	-1	95	volvo	0	1	4.0	1	0	109.1	188.8	...	0	1	0
202	-1	95	volvo	0	0	4.0	1	0	109.1	188.8	...	0	1	0
203	-1	95	volvo	1	1	4.0	1	0	109.1	188.8	...	0	1	0
204	-1	95	volvo	0	1	4.0	1	0	109.1	188.8	...	0	1	0

205 rows × 36 columns

```
# 3.17 fuel-system
```

```
df = pd.get_dummies(df, columns = ['fuel-system'] )
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	...	engine-type_ohcv	engine-type_rotor	fuel-system_1bbl :
0	3	115	alfa-romero	0	0	2.0	1	0	88.6	168.8	...	0	0	0
1	3	115	alfa-romero	0	0	2.0	1	0	88.6	168.8	...	0	0	0
2	1	115	alfa-romero	0	0	2.0	1	0	94.5	171.2	...	1	0	0
3	2	164	audi	0	0	4.0	0	0	99.8	176.6	...	0	0	0
4	2	164	audi	0	0	4.0	2	0	99.4	176.6	...	0	0	0
...
200	-1	95	volvo	0	0	4.0	1	0	109.1	188.8	...	0	0	0
201	-1	95	volvo	0	1	4.0	1	0	109.1	188.8	...	0	0	0
202	-1	95	volvo	0	0	4.0	1	0	109.1	188.8	...	1	0	0
203	-1	95	volvo	1	1	4.0	1	0	109.1	188.8	...	0	0	0
204	-1	95	volvo	0	1	4.0	1	0	109.1	188.8	...	0	0	0

205 rows × 43 columns

df.T

	0	1	2	3	4	5	6	7	8	9	...	195	196	197	198
symboling	3	3	1	2	2	2	1	1	1	0	...	-1	-2	-1	-2
normalized-losses	115	115	115	164	164	115	158	115	158	115	...	74	103	74	103
make	alfa-romero	alfa-romero	alfa-romero	audi	...	volvo	volvo	volvo	volvo						
fuel-type	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0
aspiration	0	0	0	0	0	0	0	0	1	1	...	0	0	0	1
num-of-doors	2.0	2.0	2.0	4.0	4.0	2.0	4.0	4.0	4.0	2.0	...	4.0	4.0	4.0	4.0
drive-wheels	1	1	1	0	2	0	0	0	0	2	...	1	1	1	1

3.18 bore

wheel-base	88.6	88.6	94.5	99.8	99.4	99.8	105.8	105.8	105.8	99.5	...	104.3	104.3	104.3	104.3
<code>df["bore"] = df["bore"].fillna(df["bore"].median()).astype(float)</code>															
<code>df['stroke'] = df['stroke'].fillna(df['bore'].median()).astype(float)</code>															
<code>df['horsepower'] = df['horsepower'].fillna(df['bore'].median()).astype(float)</code>															
<code>df['peak-rpm'] = df['peak-rpm'].fillna(df['peak-rpm'].median()).astype(float)</code>															
<code>df['price'] = df['price'].fillna(df['price'].median()).astype(float)</code>															

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 43 columns):
 #   Column          Non-Null Count  Dtype  
 --- 
 0   symboling       205 non-null    int64  
 1   normalized-losses 205 non-null    int64  
 2   make            205 non-null    object  
 3   fuel-type        205 non-null    int64  
 4   aspiration       205 non-null    int64  
 5   num-of-doors     205 non-null    float64 
 6   drive-wheels     205 non-null    int64  
 7   engine-location   205 non-null    int64  
 8   wheel-base        205 non-null    float64 
 9   length           205 non-null    float64 
 10  width            205 non-null    float64 
 11  height           205 non-null    float64 
 12  curb-weight      205 non-null    int64  
 13  num-of-cylinders 205 non-null    int64  
 14  engine-size       205 non-null    int64  
 15  bore             205 non-null    float64 
 16  stroke           205 non-null    float64 
 17  compression-ratio 205 non-null    float64 
 18  horsepower        205 non-null    float64 
 19  peak-rpm          205 non-null    float64 
 20  city-mpg          205 non-null    int64  
 21  highway-mpg        205 non-null    int64  
 22  price             205 non-null    float64 
 23  body-style_convertible 205 non-null    uint8  
 24  body-style_hardtop 205 non-null    uint8  
 25  body-style_hatchback 205 non-null    uint8  
 26  body-style_sedan   205 non-null    uint8  
 27  body-style_wagon   205 non-null    uint8  
 28  engine-type_dohc  205 non-null    uint8  
 29  engine-type_dohcv 205 non-null    uint8  
 30  engine-type_l      205 non-null    uint8  
 31  engine-type_ohc   205 non-null    uint8  
 32  engine-type_ohcf  205 non-null    uint8  
 33  engine-type_ocvc  205 non-null    uint8  
 34  engine-type_rotor 205 non-null    uint8  
 35  fuel-system_1bbl  205 non-null    uint8  
 36  fuel-system_2bbl  205 non-null    uint8  
 37  fuel-system_4bbl  205 non-null    uint8  
 38  fuel-system_idi   205 non-null    uint8  
 39  fuel-system_mfi   205 non-null    uint8  
 40  fuel-system_mpfi  205 non-null    uint8  
 41  fuel-system_spdi  205 non-null    uint8  
 42  fuel-system_spfi  205 non-null    uint8  
dtypes: float64(11), int64(11), object(1), uint8(20)
memory usage: 41.0+ KB
```

fuel-system_spfi	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0
-------------------------	---	---	---	---	---	---	---	---	---	---	-----	---	---	---	---

4 . Fetaure Selection

`df.corr()`


```
<ipython-input-66-2f6f6606aa2c>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version
df.corr()

symboling normalized-losses fuel-type aspiration num-of-doors drive-wheels engine-location wheel-base length width ... engine-type_ohc
symboling 1.000000 0.457484 -0.194311 -0.059866 -0.663595 -0.111150 0.212471 -0.531954 -0.357612 -0.232919 ... -0.013591
normalized-losses 0.457484 1.000000 -0.104668 -0.011273 -0.348850 0.133824 -0.021510 -0.073709 -0.006837 0.058378 ... 0.130711
fuel-type -0.194311 -0.104668 1.000000 0.401397 0.188496 0.051874 -0.040070 0.308346 0.212679 0.233880 ... -0.085551
aspiration -0.059866 -0.011273 0.401397 1.000000 0.052803 0.153897 -0.057191 0.257611 0.234539 0.300567 ... -0.070071

# 5. Train Test Split

drive-wheels -0.111150 0.133824 0.051874 0.153897 -0.003230 1.000000 0.113823 0.366828 0.416076 0.376554 ... 0.139451

df = df.select_dtypes(exclude = object) # we drop column having Data type is " object "

df.info() # We fill all NaN Values & Only int and float data type available....Now we ready for further operations

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 42 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   symboling        205 non-null    int64  
 1   normalized-losses 205 non-null    int64  
 2   fuel-type         205 non-null    int64  
 3   aspiration        205 non-null    int64  
 4   num-of-doors      205 non-null    float64
 5   drive-wheels      205 non-null    int64  
 6   engine-location    205 non-null    int64  
 7   wheel-base         205 non-null    float64
 8   length             205 non-null    float64
 9   width              205 non-null    float64
 10  height             205 non-null    float64
 11  curb-weight        205 non-null    int64  
 12  num-of-cylinders   205 non-null    int64  
 13  engine-size        205 non-null    int64  
 14  bore               205 non-null    float64
 15  stroke             205 non-null    float64
 16  compression-ratio  205 non-null    float64
 17  horsepower          205 non-null    float64
 18  peak-rpm            205 non-null    float64
 19  city-mpg            205 non-null    int64  
 20  highway-mpg          205 non-null    int64  
 21  price               205 non-null    float64
 22  body-style_convertible 205 non-null    uint8  
 23  body-style_hardtop   205 non-null    uint8  
 24  body-style_hatchback 205 non-null    uint8  
 25  body-style_sedan     205 non-null    uint8  
 26  body-style_wagon     205 non-null    uint8  
 27  engine-type_dohc     205 non-null    uint8  
 28  engine-type_dohcv    205 non-null    uint8  
 29  engine-type_l         205 non-null    uint8  
 30  engine-type_ohc       205 non-null    uint8  
 31  engine-type_ohcf      205 non-null    uint8  
 32  engine-type_ohcv      205 non-null    uint8  
 33  engine-type_rotor     205 non-null    uint8  
 34  fuel-system_1bbl      205 non-null    uint8  
 35  fuel-system_2bbl      205 non-null    uint8  
 36  fuel-system_4bbl      205 non-null    uint8  
 37  fuel-system_idi        205 non-null    uint8  
 38  fuel-system_mfi        205 non-null    uint8  
 39  fuel-system_mpfi       205 non-null    uint8  
 40  fuel-system_spdi       205 non-null    uint8  
 41  fuel-system_spfi       205 non-null    uint8  
dtypes: float64(11), int64(11), uint8(20)
memory usage: 39.4 KB

tne ohcf 0.057510 -0.210771 -0.052304 -0.054450 0.019551 0.197007 0.455127 -0.105150 -0.110320 -0.124440 ... -0.075171

x = df.drop("price",axis = 1) # It returns all Independent Variables
y = df["price"] # It returns all Dependent Variables

train, test, y_train, y_test = train_test_split(x,y,test_size=0.2,random_state=34)
x_train
```

Prediction of selling price of used cars.ipynb - Colaboratory

symboling	normalized-losses	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	width	...	engine-type_ohcv	engine-type_rotor	fuel-system_1bbl	s
203	-1	95	1	1	4.0	1	0	109.1	188.8	68.9	...	0	0	0
204	-1	95	0	1	4.0	1	0	109.1	188.8	68.9	...	0	0	0
52	1	104	0	0	2.0	0	0	93.1	159.1	64.2	...	0	0	0
2	1	115	0	0	2.0	1	0	94.5	171.2	65.5	...	1	0	0
46	2	115	0	0	2.0	1	0	96.0	172.6	65.2	...	0	0	0
...
68	-1	93	1	1	4.0	1	0	110.0	190.9	70.3	...	0	0	0
85	1	125	0	0	4.0	0	0	96.3	172.4	65.4	...	0	0	0
105	3	194	0	1	2.0	1	0	91.3	170.7	67.9	...	1	0	0
122	1	154	0	0	4.0	0	0	93.7	167.3	63.8	...	0	0	0
161	0	91	0	0	4.0	0	0	95.7	166.3	64.4	...	0	0	0

164 rows × 41 columns

x_train.shape

(164, 41)

6. Model Training

```
model = LinearRegression()
model.fit(x_train,y_train)
```

```
▼ LinearRegression
LinearRegression()
```

7. Model Evaluation

y_pred_test = model.predict(x_test) # To find Predicted Values

#Testing data evalution

y_pred = model.predict(x_test)
y_pred[20:25]

```
array([11100.00500014, 19909.19320479, 10866.041479 , 7933.8058285 ,
7669.28848162])
```

y_pred_test[20:25] # Predicted values

```
array([11100.00500014, 19909.19320479, 10866.041479 , 7933.8058285 ,
7669.28848162])
```

y_test[20:25] # actual values

```
123      8921.0
178     16558.0
3       13950.0
76      5389.0
79      7689.0
Name: price, dtype: float64
```

8. Testing Data Evalution

```
mse = mean_squared_error(y_test,y_pred_test)
print("MSE is:",mse)
rmse = np.sqrt(mse)
print("RMSE:",rmse)
mae = mean_absolute_error(y_test,y_pred_test)
print("MAE is:",mae)
r2 = r2_score(y_test,y_pred_test)
```

```
print("r2 score is :",r2)
```

```
MSE is: 14154182.910287894
RMSE: 3762.204527971319
MAE is: 2392.8348790513933
r2 score is : 0.7301625587922318
```

```
# 9. Training Data Evaluation
y_pred_train = model.predict(x_train)
mse = mean_squared_error(y_train,y_pred_train)
print("MSE is:",mse)
rmse = np.sqrt(mse)
print("RMSE:",rmse)
mae = mean_absolute_error(y_train,y_pred_train)
print("MAE is:",mae)
r2 = r2_score(y_train,y_pred_train)
print("r2 score is :",r2)
```

```
MSE is: 4974235.837890235
RMSE: 2230.299495110519
MAE is: 1585.3408437358958
r2 score is : 0.9223636932435281
```